

Quick Intro: Attention Heads and Transformer Layers

Paul Rodriguez PhD

2023

Outline

- Basic word prediction task and motivating the attention strategy
- A basic Attention Head network
- Transformers
- Exercise: Working with BERT pretrained transformer model

Dependences of Language

Consider this sequence:

*The Law will never be perfect, but it's application
should be just - this is what we are missing, in my
opinion <End of Sequence>*

What does 'it' refer to that can have an 'application'?

Dependences of Language


Consider this sequence:

The Law will never be perfect, but it's application should be just - this is what we are missing, in my opinion <End of Sequence>

What does 'it' refer to that can have an 'application'?

e.g 'it' refers back to 'Law', which is part of 'the Law' noun phrase, which is the entity that will 'never be perfect', and so on ...

many dependencies and interactions



A toy problem to get some intuition

- Let's use the following list of 5 tokens:
 <start>, the, man, chicken, ordered
- Let's use this sequence of 6 tokens as our only:
 <start> the man ordered the chicken
- If we use token ids 1 to 5 it is the sequence of 6 numbers [1,2,3,5,2,4]
- **Now let's try to predict the next word by 'attention' idea**

The toy task: predict next word

The data: 5 tokens ($V=5$),

1 sequence (length= $T=6$): <Start> the man ordered the chicken

A basic solution is bigram matrix

eg a sequence of tokens (rows) and current word predictions (cols)

$X = \text{Sequence-to-Word is } T \times V$

Pos	Word	<strt>	The	Man	Chikn	Order
0	<start>		1.0			
1	The			0.5	0.5	
2	Man					1.0
3	Orde.r		1.0			
4	The			0.5	0.5	
5	Chick.	1.0				

The toy task: predict next word

The data: 5 tokens ($V=5$),

1 sequence (length= $T=6$): <Start> the man ordered the chicken

A basic solution is bigram matrix

eg a sequence of tokens (rows) and current word predictions (cols)

X= Sequence-to-Word is $T \times V$

Challenge, can we learn predictions (\rightarrow) that depend on context of other tokens and/or position

After <Start> the \rightarrow man = 1.0

After 'Ordered' the \rightarrow chicken = 1.0

Pos	Word	<strt>	The	Man	Chikn	Order
0	<start>		1.0			
1	The			0.5	0.5	
2	Man					1.0
3	Orde.r		1.0			
4	The			0.5	0.5	
5	Chick.	1.0				

The attention idea

Let's get all tokens to 'pass information' about dependencies

E.G. for X a $T \times V$ matrix, we want W a $T \times T$ matrix – aka 'attention' weights - so that $W * X = T \times V$ has contextual predictive information

W dependencies is $T \times T$

X= Sequence-to-Word is $T \times V$

$$\begin{pmatrix} w_{11} & \cdots & w_{1T} \\ \vdots & \vdots & \vdots \\ w_{T1} & \cdots & w_{TT} \end{pmatrix} * \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 1.0 & 0 & 0 & 0 & 0 \end{pmatrix} \rightarrow \rightarrow \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 \end{pmatrix}$$

The attention idea

Let's get all tokens to 'pass information' about dependencies

E.G. for X a $T \times V$ matrix, we want W a $T \times T$ matrix – aka 'attention' weights - so that $W * X = T \times V$ has contextual predictive information

W dependencies is $T \times T$

X= Sequence-to-Word is $T \times V$

$$\begin{pmatrix} w_{11} & \cdots & w_{1T} \\ \vdots & \vdots & \vdots \\ w_{T1} & \cdots & w_{TT} \end{pmatrix} * \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{pmatrix} \rightarrow \rightarrow \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 \end{pmatrix}$$

W will be learned and should depend on transformations of X

Making predictions causal

(ie only depends on previous tokens)

First build a TxT mask so that sequence position t only uses columns $1:t$

$$Mask = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 & 0 \\ 1 & 1 & 1 & 1 & 1 & 1 \end{pmatrix}$$

Then, multiply it by W_{dep} matrix

Now multiply masked dependency elementwise to W_{dep}


Finally, apply Softmax to each row to make it like probability weights

$$W_{dep} \odot Mask = \begin{pmatrix} w_{11} & 0 & \cdots & \cdots & 0 \\ w_{21} & w_{22} & 0 & \cdots & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ w_{T1} & w_{T2} & w_{T3} & \cdots & w_{TT} \end{pmatrix}$$


Transformation as Embedding Layer/Table

Let X be given as a 'one-hot vector input' – where one of the 25 input nodes is 1, the rest are 0.

Use linear activation so that $Y = W * X$



X is 25x1




Transformation as Embedding Layer/Table

Let X be given as a 'one-hot vector input' – where one of the 25 input nodes is 1, the rest are 0.


Use linear activation so that $Y = W * X$

If j -th node is 1, then for W a 10×25 matrix

$$Y = W * X$$
$$= \begin{pmatrix} | & | & \cdot & | & \cdot & \cdot \\ w_1 & w_2 & \cdot & w_j & \cdot & \cdot \\ | & | & \cdot & | & \cdot & \cdot \end{pmatrix} * \begin{pmatrix} 0 \\ 0 \\ \cdot \\ 1 \\ \cdot \\ \cdot \end{pmatrix} = ?$$



X is 25x1



Transformation as Embedding Layer/Table

Let X be given as a 'one-hot vector input' – where one of the 25 input nodes is 1, the rest are 0.

Use linear activation so that $Y = W * X$


if j -th node is 1, then for W a 10×25 matrix

$$Y = W * X = W_j = \begin{pmatrix} | \\ w_j \\ | \end{pmatrix}$$

So just let X be sequence of token ids and treat W like a table

X is 25x1

Y is 10x1



An Attention Head construction

1. Let X = sequence of token ids

For embedding dimension E , map:

$$\underset{T \times 1}{X} \rightarrow \underset{T \times E}{X_{emb}}$$

For positions $1 \dots T$ map:

$$\underset{T \times 1}{P} \rightarrow \underset{T \times E}{X_{pos}}$$

Take as new input:

$$X = X_{emb} + X_{pos}$$

2. For dimension H , get Q, K, V matrices: $\underset{T \times E}{X} \rightarrow \underset{T \times H}{Query}$ $\underset{T \times E}{X} \rightarrow \underset{T \times H}{Key}$ $\underset{T \times E}{X} \rightarrow \underset{T \times H}{Value}$

3. Get Attention Weights and Output:

$$\underset{T \times H}{Q} * \underset{H \times T}{K'} \rightarrow \underset{T \times T}{W_{dep}}$$

$$Output = softmax(W_{dep} \odot Mask) * V$$

All mappings are
linear transformations



An Attention Head code

Hyperparameters:

E is for embedding dimension,

H is for “head size” (dimension) of transformations (we could let $E = H$)

11 Layers/Functions

```
#Now build model to learn transformation for Q,K,V matrices

Xsequence      = tf.keras.layers.Input(shape=(T,V)) #the batch size is left unspecified
Pos_Input      = tf.keras.layers.Input(shape=(T)) #just the t=1...T integer
Pos_Embed      = tf.keras.layers.Embedding(T,V, input_length=T,name='PosEmbed')(Pos_Input) #input wi
Xinputs        = tf.keras.layers.Add()([Xsequence, Pos_Embed])

#now feed to Q,K,V transformations
Qmat           = tf.keras.layers.Dense(H,activation='linear',use_bias=False,name='Qmat')(Xinputs) #so f
Kmat           = tf.keras.layers.Dense(H,activation='linear',use_bias=False,name='Kmat')(Xinputs)
Vmat           = tf.keras.layers.Dense(H,activation='linear',use_bias=False,name='Vmat')(Xinputs)

#now apply QtoK take softmax, scale it , apply to V
QK             = tf.keras.layers.Dot(axes=(2))([Qmat,Kmat]) #it will treat each Batch item separately
QKscaled       = tf.keras.layers.Lambda(lambda x: x * scale_constant)(QK) #for each x in QK mult by
Attn_Wts       = tf.keras.layers.Softmax(axis=2,name='AttnWts')(QKscaled, mask=Msk) #apply mas
Vout           = tf.keras.layers.Dot(axes=1,name='Voutput')([Attn_Wts,Vmat])

my_attn_model  = tf.keras.Model(inputs = [Xsequence,Pos_Input], outputs=Vout)
```


An Attention Head code

Hyperparameters:

E is for embedding dimension,

H is for “head size” (dimension) of transformations (we could let $E = H$)

11 Layers/Functions

```
#Now build model to learn transformation for Q,K,V matrices


Xsequence      = tf.keras.layers.Input(shape=(T,V)) #the batch size is left unspecified
Pos_Input      = tf.keras.layers.Input(shape=(T)) #just the t=1...T integer
Pos_Embed      = tf.keras.layers.Embedding(T,V, input_length=T,name='PosEmbed')(Pos_Input) #input wi
Xinputs        = tf.keras.layers.Add()([Xsequence, Pos_Embed])

#now feed to Q,K,V transformations
Qmat           = tf.keras.layers.Dense(H,activation='linear',use_bias=False,name='Qmat')(Xinputs) #so f
Kmat           = tf.keras.layers.Dense(H,activation='linear',use_bias=False,name='Kmat')(Xinputs)
Vmat           = tf.keras.layers.Dense(H,activation='linear',use_bias=False,name='Vmat')(Xinputs)

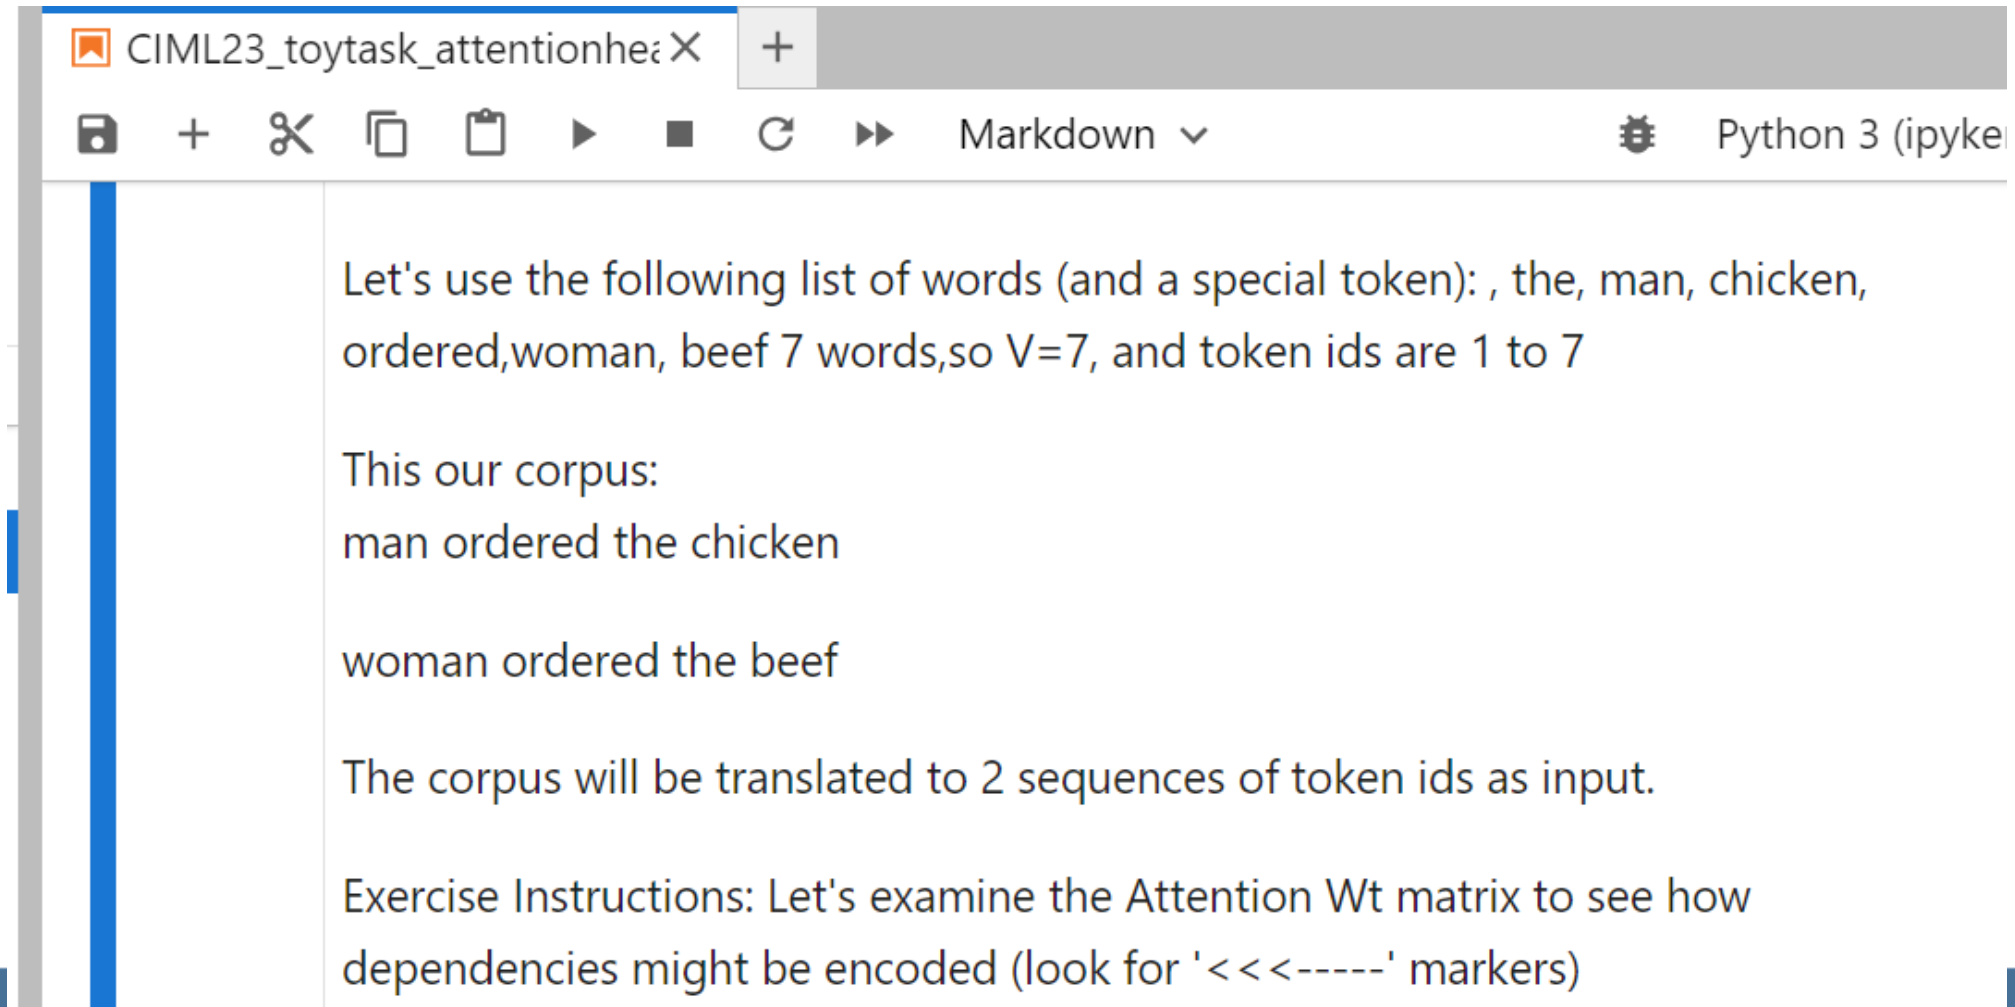
#now apply QtoK take softmax, scale it , apply to V
QK             = tf.keras.layers.Dot(axes=(2))([Qmat,Kmat]) #it will treat each Batch item separately
QKscaled       = tf.keras.layers.Lambda(lambda x: x * scale_constant)(QK) #for each x in QK mult by
Attn_Wts       = tf.keras.layers.Softmax(axis=2,name='AttnWts')(QKscaled, mask=Msk) #apply mas
Vout           = tf.keras.layers.Dot(axes=(1,name='Voutput'))([Attn_Wts,Vmat])

my_attn_model  = tf.keras.Model(inputs = [Xsequence,Pos_Input], outputs=Vout)
```

Also a scaling by H

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$


An example of attention head with a toy task:
(in the CIML23_toytask_attention notebook)



The screenshot shows a Jupyter Notebook window titled "CIML23_toytask_attentionhe". The toolbar includes icons for saving, adding, deleting, copying, pasting, running, and a dropdown menu set to "Markdown". The notebook content is as follows:

Let's use the following list of words (and a special token): , the, man, chicken, ordered,woman, beef 7 words,so $V=7$, and token ids are 1 to 7

This our corpus:

man ordered the chicken


woman ordered the beef

The corpus will be translated to 2 sequences of token ids as input.

Exercise Instructions: Let's examine the Attention W_t matrix to see how dependencies might be encoded (look for '<<<-----' markers)

Output TxV predictions:

Notice that the→ [chicken or beef]
predictions change



1/1 [=====] - 0s 106ms/step

This is the output predictions use Head size 20

	<ST>	the	man	chkn	ordrd	woman	beef
0 <ST>	0.596	0.037	0.988	0.202	0.193	0.988	0.077
1 man	0.062	0.579	0.575	0.243	0.996	0.184	0.802
2 ordrd	0.092	0.991	0.122	0.612	0.319	0.050	0.403
3 the	0.225	0.705	0.016	0.996	0.008	0.018	0.837
4 chkn	0.994	0.103	0.365	0.679	0.104	0.634	0.553

This is the output predictions use Head size 20

	<ST>	the	man	chkn	ordrd	woman	beef
0 <ST>	0.596	0.037	0.988	0.202	0.193	0.988	0.077
1 woman	0.037	0.578	0.189	0.061	0.999	0.051	0.940
2 ordrd	0.129	0.994	0.049	0.626	0.153	0.034	0.511
3 the	0.028	0.686	0.163	0.727	0.786	0.034	0.992
4 beef	0.994	0.107	0.369	0.679	0.099	0.633	0.340

Output TxV predictions:

Notice that the → [chicken or beef] predictions change

1/1 [=====] - 0s 106ms/step

This is the output predictions use Head size 20

	<ST>	the	man	chkn	ordrd	woman	beef
0 <ST>	0.596	0.037	0.988	0.202	0.193	0.988	0.077
1 man	0.062	0.579	0.575	0.243	0.996	0.184	0.802
2 ordrd	0.092	0.991	0.122	0.612	0.319	0.050	0.403
3 the	0.225	0.705	0.016	0.996	0.008	0.018	0.837
4 chkn	0.994	0.103	0.365	0.679	0.104	0.634	0.553

This is the output predictions use Head size 20

	<ST>	the	man	chkn	ordrd	woman	beef
0 <ST>	0.596	0.037	0.988	0.202	0.193	0.988	0.077
1 woman	0.037	0.578	0.189	0.061	0.999	0.051	0.940
2 ordrd	0.129	0.994	0.049	0.626	0.153	0.034	0.511
3 the	0.028	0.686	0.163	0.727	0.786	0.034	0.992
4 beef	0.994	0.107	0.369	0.679	0.099	0.633	0.340

TxT Attn Wts:

Notice that the ← wts are from [ordered or man]

for i-th input: 0 These are the output at layer AttnWts

	0 <ST>	1 man	2 ordrd	3 the	4 chkn
0 <ST>	1.000	0.000	0.000	0.000	0.000
1 man	0.093	0.907	0.000	0.000	0.000
2 ordrd	0.064	0.561	0.375	0.000	0.000
3 the	0.003	0.213	0.759	0.025	0.000
4 chkn	0.258	0.003	0.160	0.523	0.056

The head size H was: 20

for i-th input: 1 These are the output at layer AttnWts


	0 <ST>	1 woman	2 ordrd	3 the	4 beef
0 <ST>	1.000	0.000	0.000	0.000	0.000
1 woman	0.079	0.921	0.000	0.000	0.000
2 ordrd	0.091	0.374	0.535	0.000	0.000
3 the	0.001	0.769	0.223	0.007	0.000
4 beef	0.252	0.001	0.162	0.526	0.059

The head size H was: 20

Finally, a transformer

Include skip-add connections

Include Layer Normalization or DropOut layers

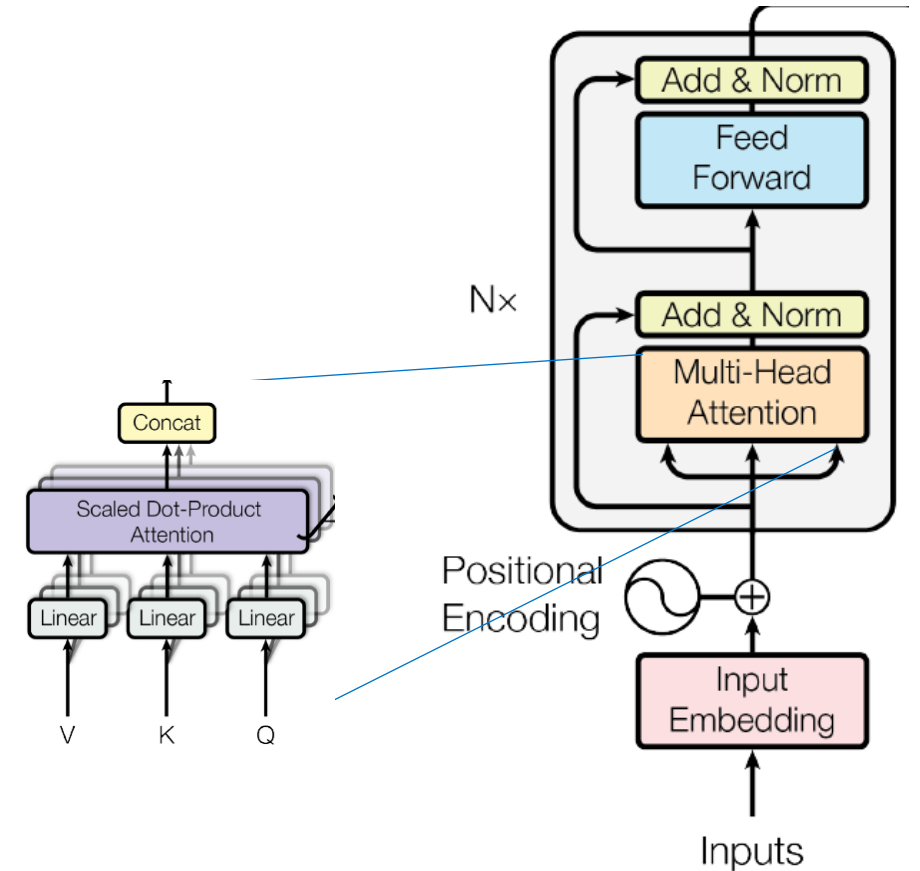


Finally, a transformer

Include skip-add connections

Include Layer Normalization or DropOut layers

Multi-Head – for N heads produce $T \times (H/N)$ each




Finally, a transformer

Include skip-add connections

Include Layer Normalization or DropOut layers

Multi-Head – for N heads produce $T \times (H/N)$ each

Add MLP layers on top –
output another $T \times E$ matrix **stackable!**



Finally, a transformer

Include skip-add connections

Include Layer Normalization or DropOut layers

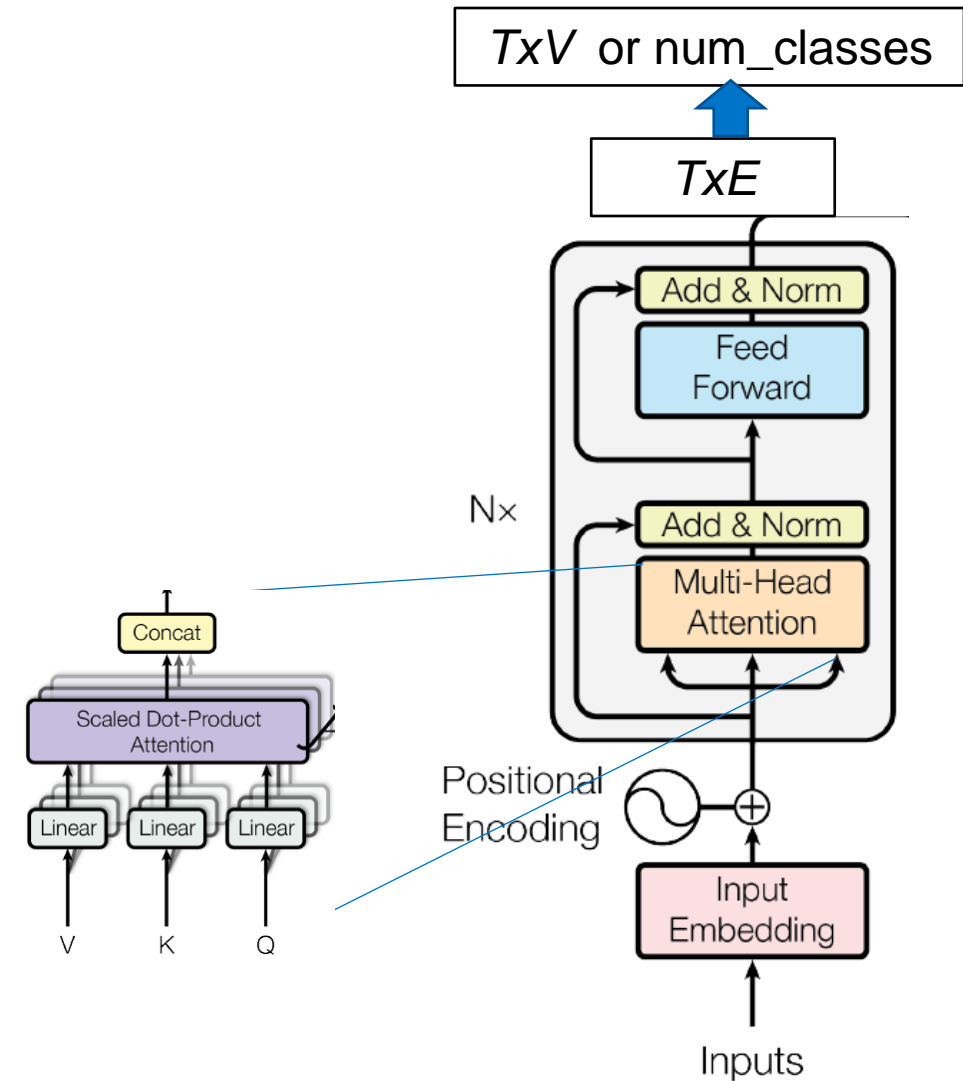
Multi-Head – for N heads produce $T \times (H/N)$ each

Add MLP layers on top –

output another $T \times E$ matrix

stackable!

or output final probabilities




2 kinds of training strategies

GPT – predict next word only look back at prior context
(which could be a whole document)

Put mask on attention weights so that predictions only depend on previous tokens

BERT – ***No attention mask*** so all token dependencies can influence all other tokens predictions

Special tokens help create a variety of tasks

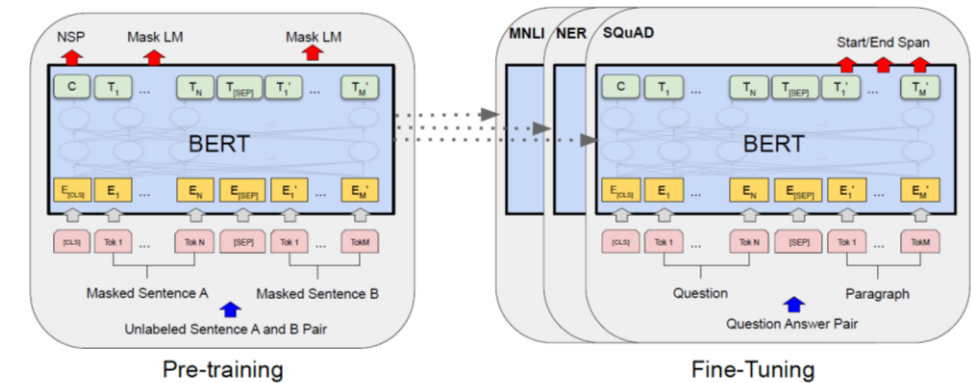


BERT (Bidirectional Encoder Representations from Transformers)

Goal: Train a model to develop general token-level AND sentence-level encoding

1 Pretrain on:

- fill-in-the-blank
- binary classification if 2 sentences go together



Devlin, etal, 2019

2 Fine tune on variety of tasks


BERT Input: 2 sentences

[CLS]
is 'dummy' token
and always
present

[SEP] marks separation
between sentences or
sequences

[PAD] tokens help
fill out sequence to T
items

Input



Devlin, etal, 2019

BERT Input: 2 sentences

[CLS]
is 'dummy' token
and always
present

[SEP] marks separation
between sentences or
sequences

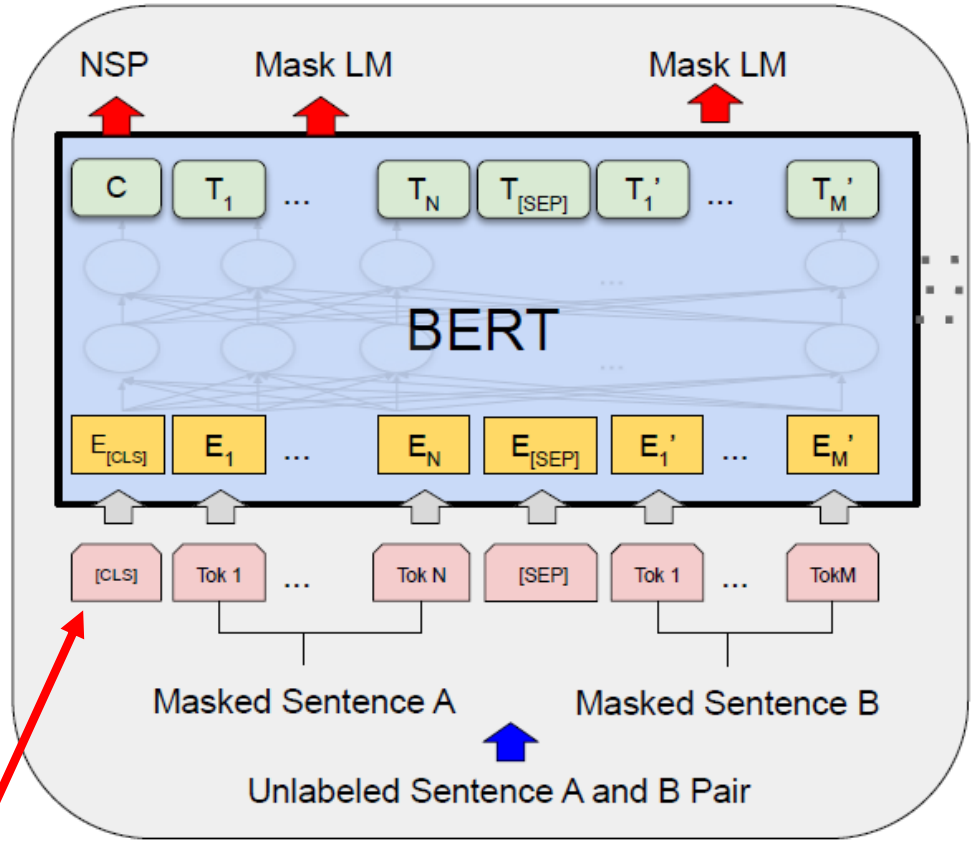
[PAD] tokens help
fill out sequence to T
items

Input	[CLS]	my	dog	is	cute	[SEP]	he	likes	play	##ing	[SEP]
Token Embeddings	$E_{[CLS]}$	E_{my}	E_{dog}	E_{is}	E_{cute}	$E_{[SEP]}$	E_{he}	E_{likes}	E_{play}	$E_{##ing}$	$E_{[SEP]}$
	+	+	+	+	+	+	+	+	+	+	+
Segment Embeddings	E_A	E_A	E_A	E_A	E_A	E_A	E_B	E_B	E_B	E_B	E_B
	+	+	+	+	+	+	+	+	+	+	+
Position Embeddings	E_0	E_1	E_2	E_3	E_4	E_5	E_6	E_7	E_8	E_9	E_{10}

Add a TxE
matrix for 'A'
vs 'B' sentence

BERT –

Last layer $T \times V$ output
Transformer Layers
Embeddings
T Input tokens



[CLS]
is 'dummy' token
and always
present

Pre-training

Devlin, etal, 2019

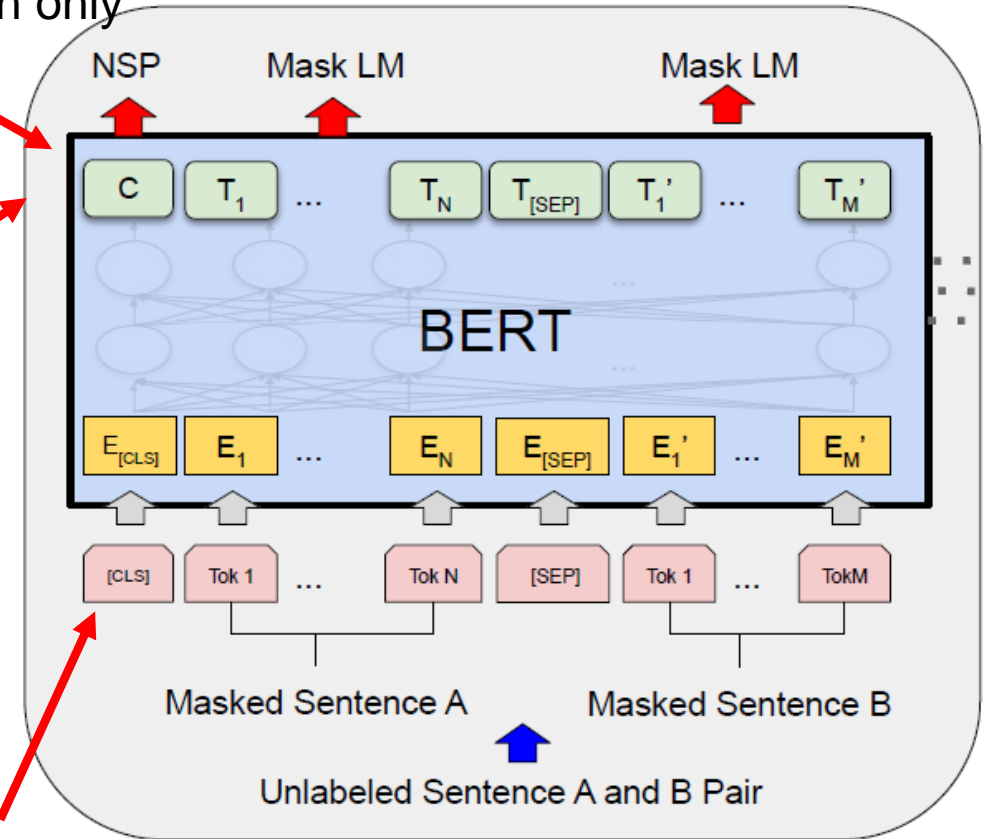
Train a classification task on this
[CLS] token prediction only

BERT –

For classification over paired sentences or any text, train C row vector (in the last layer $T \times V$ output matrix)

“NSP”: classify if next sentence follows first sentence

[CLS]
is 'dummy' token
and always
present



KerasNLP

- KerasNLP is an extension to Keras
- KerasNLP has several pre-trained LLMs (large language models). Each model comes with related modules, for example:
 - GPT2Backbone the model without task specific output layers
 - GPT2CausalLM the model with output predictions
 - GPT2CausalLMPreprocessor the preprocessor that feeds model.fit

KerasNLP

- KerasNLP is an extension to Keras
- KerasNLP has several pre-trained LLMs (large language models). Each model comes with related modules, for example:
 - GPT2Backbone the model without task specific output layers
 - GPT2CausalLM the model with output predictions
 - GPT2CausalLMPreprocessor the preprocessor that feeds model.fit

We will use pre-trained BERT and compare different BERT versions

Keras NLP package is not in the tensorflow container, so we can do:

1. Rebuild Container - check
2. Set up conda environment
3. Run 'pip install' commands

Option 3 is quick and a bit dirtier, but often useful and part of tutorials

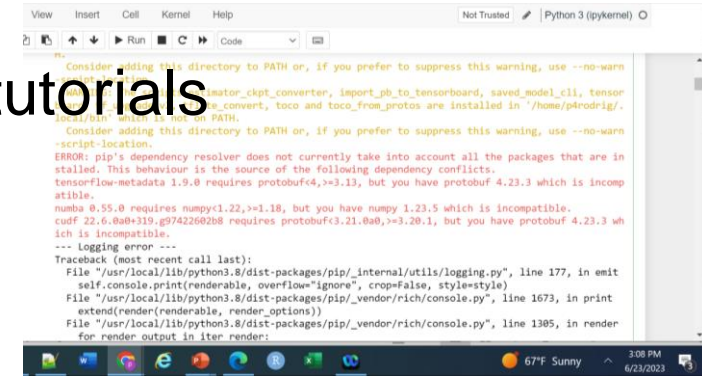
On Expanse you might get warnings or logging errors.

But it should show the package as installed

Note, 'pip install' will put packages in .local directory which is on PythonPath. So there's a potential to clash with other notebooks – you may have to move/manage your .local directories

```
In [6]: !pip install -q --upgrade keras-nlp
```

```
WARNING: The scripts f2py, f2py3 and f2py3.8 are installed in '/home/p4rodrig/.local/bin' which is not on PATH
```



```
In [7]: !pip show keras_nlp
```

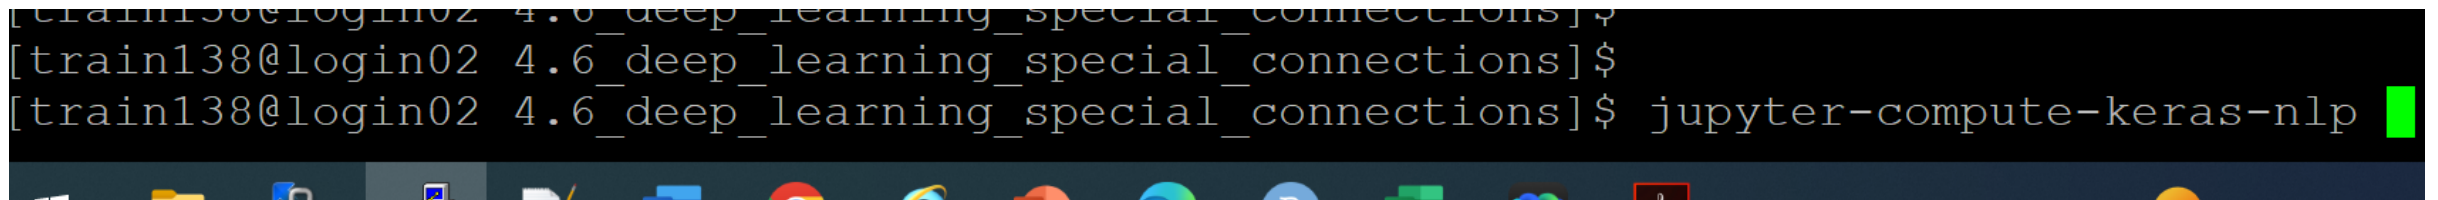
```
Name: keras-nlp
Version: 0.5.2
Summary: Industry-strength Natural Language Processing extensions for Keras.
Home-page: https://github.com/keras-team/keras-nlp
Author: Keras team
Author-email: keras-nlp@google.com
License: Apache License 2.0
Location: /home/p4rodrig/.local/lib/python3.8/site-packages
Requires: absl-py, numpy, packaging, tensorflow-text
Required-by:
```

```
In [8]: !import os
```


Notebook exercise using KerasNLP

- In a terminal window start the notebook session for keras-nlp

```
...]$ cd 4.6_deep_learning_special_connections
...]$ jupyter-compute-keras-nlp
```

A screenshot of a terminal window with a black background and white text. The prompt is [train138@login02 4.6_deep_learning_special_connections]\$ and the command being entered is jupyter-compute-keras-nlp. A green cursor is visible at the end of the command. The terminal window is part of a desktop environment with a taskbar at the bottom showing various application icons.

- Open the URL and look for the *CIML23_BERT_FineTune_v3.ipynb* notebook

Keras NLP package has several BERT versions

So let's start with 'bert-small' (28M parameters)

For reference:

BERT_{BASE} (L=12, H=768, Attn=12, Total Parameters=110M)

BERT_{LARGE} (L=24, H=1024, Attn=16, Total Parameters=340M).



models/			
bert_tiny_en_uncased	BERT	4M	2-layer BERT model where all input is lowercased. Trained on English Wikipedia + BooksCorpus.
bert_small_en_uncased	BERT	28M	4-layer BERT model where all input is lowercased. Trained on English Wikipedia + BooksCorpus.
bert_medium_en_uncased	BERT	41M	8-layer BERT model where all input is lowercased. Trained on English Wikipedia + BooksCorpus.
bert_base_en_uncased	BERT	109M	12-layer BERT model where all input is lowercased. Trained on English Wikipedia + BooksCorpus.
bert_base_en	BERT	108M	12-layer BERT model where case is maintained. Trained on English Wikipedia + BooksCorpus.
bert_base_zh	BERT	102M	12-layer BERT model. Trained on Chinese Wikipedia.
bert_base_multi	BERT	177M	12-layer BERT model where case is maintained. Trained on trained on Wikipedias of 104 languages

END